

Idea for a project: KR for ML

December 19, 2016

Abstract

Hypothesis: ML has focuses on solving concrete learning/prediction tasks. In order to harness this technology, it needs to be imbued with a KR philosophy.

1 Next steps

1. What is the killer-app? can combination of ontology and ML help guide ML techniques? what does one need ontologies to do? <http://dl-learner.org/about/ontology-learning/>
2. <http://www.ttic.edu/SNL2017/>
3. Find and engage the best ML people interested in CAV/KR.
4. Do a literature search to find out what has been done at the interface of ML and KR (see Section 5).
 - (a) <https://sites.google.com/site/krr2015/home/schedule> was a workshop on symbolic and neural approaches to KRR.
 - (b) <http://www1.icsi.berkeley.edu/~shastri/shruti/> has work on modeling logic programs as NN and running the NN to evaluate queries.
 - (c) <http://www.ntu.edu.sg/home/asahtan/> has worked on integrating domain knowledge into NNs.
 - (d) Lise Getoor: "we need machine learning for graphs"
5. Can one solve synthesis using ML? Yes, genetic programming.
6. The five tribes of ML: <https://www.youtube.com/watch?v=B8J4uefCQM>
7. https://en.wikipedia.org/wiki/Markov_logic_network
8. Pour résoudre des jeux, j'avais parlé de Daniel Neider : <https://www.lii.rwth-aachen.de/en/18-mitarbeiter/wissenschaftliche-mitarbeiter/team-prof-grohe/121-publications-of-daniel-neider.html> Celui-la par ex : <https://arxiv.org/abs/1507.05612>
Sur l'application de l'algorithme de Angluin dans la vérification compositionnelle : Cobleigh, J.M., Giannakopoulou, D., and Pasareanu, C.S. "Learning Assumptions for Compositional Verification", TACAS 2003. Voir ici : <https://ti.arc.nasa.gov/profile/dimitra/>

P Madhusudan fait des choses aussi avec D Neider entre autres : <http://madhu.cs.illinois.edu/pub.html>

9. <http://ti.tuwien.ac.at/cps/teaching/practicals/?id=178> Neurons controlling locomotion.

Glossary

- NN = artificial neural network
- KR = knowledge representation
- ML = machine learning

2 NNs 101

An *activation function* is a function $F : \mathbb{R} \rightarrow \mathbb{R}$. Two common activation functions are

$$F_1(z) = \begin{cases} 1 & \text{if } z > 0 \\ 0 & \text{otherwise.} \end{cases}$$

and

$$F_2(z) = \frac{1}{1 + e^{-z}}.$$

A *neuron* is a function $G : \mathbb{B}^n \rightarrow \mathbb{B}$ where $G(\bar{x}) = F(\bar{x} \cdot \bar{w})$ for some *weights* $w_i \in \mathbb{R}$. Neurons with activation functions F_1 are called *perceptrons*, and with activation functions F_2 are called *sigmoid neurons*.

A *neural network (NN)* is a DAG whose nodes are neurons. A NN with N inputs and M outputs computes a function $F : \mathbb{B}^N \rightarrow \mathbb{B}^M$. One may round the output of F to the nearest integer 0 or 1 to get a function $F : \mathbb{B}^N \rightarrow \mathbb{B}^M$.

Fact 1. *NN built from perceptrons can compute arbitrary Boolean functions. E.g., $N = 3, M = 1, w_1 = w_2 = 1, w_3 = -1$ determines $x_1 \wedge x_2$ if we fix $x_3 = 1$; and $N = 2, M = 1, w_1 = -2, w_2 = 1$ determines $\neg x_1$ if we fix $x_2 = 1$.*

One might be tempted to consider the synthesis problem is: given a function G , find a NN that computes it. The problem is that G is typically not known so one can't know if a NN actually computes G . Instead, some tuples $(x, G(x))$ are known. This data can be used as training and validation data, and there are measures of how well a NN represents G (e.g., cross-validation). Thus from a ML point of view, this is a *optimisation* problem which one can think of as minimising some error.

The typical approach in NN seems to be: a) get a bunch of data, b) take an educated guess at the DAG, c) apply an optimisation technique (e.g., gradient descent) to *learn* the weights that minimise the error.

Remark 1. *I read that there are no good learning algorithms for NNs using F_1 unless the DAG has very special form (e.g., consists of a single neuron).*

3 KR for NN

Suppose you've trained a bunch of NN to do certain tasks in the same domain. Are these NNs consistent with each other? Can one compose these NNs to do more complex tasks?

Example 1. Suppose N_1 recognises pictures of birds and N_2 recognises pictures of animals.

1. consistency: check if every input recognised by N_1 is also recognised by N_2 ?
2. composition: build a NN that recognises all animals that are not birds.

Definition 1. Given a finite directed graph (V, E) and a NN N_v for every $v \in V$. This setting is consistent if $(v, w) \in E$ implies $N_v(\bar{x}) \rightarrow N_w(\bar{x})$ for all \bar{x} (intuitively, E represents the “is a” relation).

Theorem 1. Assume we are using the activation function F_1 . Consistency is decidable.

Proof. Recall that the FO-theory of $\mathcal{R} := (\mathbb{R}, +, <)$ is decidable, a fact that can be proved using automata on infinite words. But every NN $N(\bar{x})$ can be written as a FO formula $A(\bar{x})$ over \mathcal{R} . Thus we can decide consistency, i.e., decide if $\forall \bar{x}. A(\bar{x}) \rightarrow B(\bar{x})$. \square

Actually, $A(\bar{x})$ is an existential formula, i.e., of the form $\exists \bar{y}. \phi(\bar{x}, \bar{y})$ where ϕ is a quantifier free formula in the signature of \mathcal{R} . So, consistency can be checked in EXPTIME. Can it be done in PSPACE?

What about composition? We can build a formula for the composition, e.g., $C(\bar{x}) = A(\bar{x}) \wedge \neg B(\bar{x})$. But, how does one turn C into a NN? Does one need to turn it into a NN?

Can we say anything if the activation function is F_2 ?

3.1 DL for ML

NN can be used to represent binary relations between objects, e.g., “friend” relation, “like” relation, etc.

So, the description logic FL^{-1} seems like a start for doing basic reasoning about the objects and their relationships. We describe this logic. Given a set of unary predicates C_i and binary predicates R_j , a *concept* is generated by the unary predicates, and closed under the following operations: $C_i \cap C_j$, $\{x : \forall y. R_j(x, y) \rightarrow C_i(y)\}$, and $\{x : \exists y. R_j(x, y)\}$.

In DL, one basic question is “subsumption”, i.e., $\forall x. C(x) \rightarrow D(x)$. This is what we called consistency before.

Issue: One problem with subsumption/consistency is that $\forall x$ is probably too strong. Perhaps one needs something like “almost all x”. Or, assuming that the data is labeled consistently, what is the error/probability that a data point x is classified as being in C but not in D .

4 Wild ideas

1. Can FOL deduction be implemented in a NN? See Smolensky 1988 "On the proper treatment of Connectionism".

5 Literature

5.0.1 Some basic machine terminology

5.0.2 Towell and Shavlik 1994, KR for ANN

5.0.3 Markov Logic Networks