

Sasha Rubin

Research Statement

June 2017

Overview

My work consists of developing and applying formal and logical methods to modeling and reasoning about computational systems, often involving strategic behaviour of agents. This is motivated by “explainable AI”, the need to ensure that the systems being built using, e.g., machine learning, can explain their decisions and actions to human users.

I approach these questions by extending classic techniques in formal methods, such as the use of games on graphs and reductions to automata, and applying theoretical insights to foundational problems in AI. One notable example of this convergence is my recent IJCAI17 paper [17] which is at the intersection of formal methods and automated planning in AI. I am also pursuing more speculative questions such as “What is synthesis and how should it be formalised?”.

Current Research — Formal methods for multi-agent systems

Multi-agent systems (MAS) involve multiple individual agents (these may be people, software, robots) each with their own goals. Such systems can be viewed as multi-player games, and thus notions from game-theory (e.g., strategies, knowledge, and equilibria) are used to reason about them. Agents in realistic MAS often lack information about other agents and the environment, and is often categorised in one of two ways: a) *incomplete information* and b) *imperfect information*.

a) MAS with incomplete information

Incomplete-information refers to uncertainty about the environment (i.e., the structure of the game). I have considered two sources of incomplete information for MAS.

First, the *number of agents* may not be known, or may not be bounded a priori. In a series of papers, I have contributed to a generalisation of a cornerstone paper on verification of such systems (“Reasoning about Rings”, E.A. Emerson, K.S. Namjoshi, POPL, 1995) from ring topologies to arbitrary topologies [1, 2, 8, 3]. Other work on this topic studied the relative power of standard communication-primitives assuming an unknown number of agents [10], as well as the complexity of model-checking timed systems assuming an unknown number of agents [9]. I also contributed to a book on this topic published by Morgan&Claypool in 2015 [15, 16].

Second, the agents may be operating in a *partially-known environment*. For instance, the agents may know they are in a ring, but may not know the size of the ring. I launched the application of automata theory for the verification of high-level properties of light-weight mobile agents in partially-known environments [30]. In follow-up work I explored this theme further, including finding ways to model agents on grids — the most common abstraction of 2D and 3D space [4, 5, 27].

b) MAS with imperfect information

Even if agents have certainty about the structure of the system, they may not know exactly which state the system is in. This is called imperfect information and the associated logic for reasoning about such cases are called *epistemic*. I have studied strategic-epistemic logics in a number of works, namely, with a prompt modality (thus allowing one to express that a property holds “promptly” rather than simply “eventually”) [6],

and on systems with public-actions (such as certain card games, including a hand of Poker or a round of Bridge) [13, 12]. The importance of these last works is that they give the first decidability (and sometimes optimal complexity) results for strategic reasoning about games of imperfect information in which the agents may have arbitrary observations. In contrast, following classical restrictions on the observations or information of agents, I have also shown how to extend strategy logic by epistemic operators and identified a decidable fragment in which one can express equilibria concepts [14].

Foundations of Automated Planning

Planning in AI can be viewed as the problem of finding strategies in one- or two-player graph-games. In this model vertices represent states, edges represent transitions, and the players represent the agents. I have contributed foundational work to such games. Concretely, I recently extended the classic belief-space construction for games of imperfect-information from finite arenas to infinite-arenas [19] (infinite arenas often arise in the study of MAS with incomplete information, see above). I have also used these ideas to elucidate the role of observation-projections in generalised planning problems [18, 17]. I have generalised classic results about certain games with quantitative objectives (i.e., Ehrenfeucht and J. Mycielski. Positional strategies for mean payoff games. *International Journal of Game Theory*, 8:109–113, 1979) to so-called first-cycle games, i.e., games in which play stops the moment a vertex is repeated [7].

Past Research — Algorithmic Model Theory

My early work contributed to a research program called “Algorithmic Model Theory” whose aim is to develop and extend the success of Finite Model Theory to infinite structures that can be reasoned about algorithmically.

Specifically, my PhD work pioneered the development of “automatic structures”: this is a generalisation of the regular languages from sets to mathematical objects with structure, such as graphs, arithmetics, algebras, etc. The fundamental property of automatic structures is that one can automatically answer logic-based queries about them (precisely, their first-order theory is decidable). I gave techniques for proving that structures are or are not automatic (similar to, but vastly more complicated than, pumping lemmas for regular languages), I studied the computational complexity of deciding when two automatic structures are the same (isomorphic), and I found extensions of the fundamental property, thus enriching the query language [11, 20, 21, 22, 23, 24, 25, 29]. I have also worked on extensions of automatic structures to include oracle computation [26, 28].

Vision for future research

My research trajectory is to help bridge two communities: formal-methods and AI. To that end, I plan to continue bringing formal- and logical-methods to bear on questions of importance to computer scientists and society. Notably, “explainable AI” (www.darpa.mil/program/explainable-artificial-intelligence) will require an integration of formal methods, knowledge representation, and machine-learning. Concretely, I plan to organise future editions of the “Formal Methods in Artificial Intelligence” workshop, and to expand my work to cover other central aspects of AI such as formal methods for Neural Networks.

References

- [1] Benjamin Aminof, Swen Jacobs, Ayrat Khalimov, and Sasha Rubin. Parameterized model checking of token-passing systems. In *Verification, Model Checking, and Abstract Interpretation - 15th International Conference, VMCAI 2014, San Diego, CA, USA, January 19-21, 2014, Proceedings*, pages 262–281, 2014.
- [2] Benjamin Aminof, Tomer Kotek, Sasha Rubin, Francesco Spegni, and Helmut Veith. Parameterized model checking of rendezvous systems. In *CONCUR 2014 - Concurrency Theory - 25th International Conference, CONCUR 2014, Rome, Italy, September 2-5, 2014. Proceedings*, pages 109–124, 2014.
- [3] Benjamin Aminof, Tomer Kotek, Sasha Rubin, Francesco Spegni, and Helmut Veith. Parameterized model checking of rendezvous systems. *Distributed Computing*, pages 1–36, 2017.
- [4] Benjamin Aminof, Aniello Murano, Sasha Rubin, and Florian Zuleger. Verification of asynchronous mobile-robots in partially-known environments. In *PRIMA 2015: Principles and Practice of Multi-Agent Systems - 18th International Conference, Bertinoro, Italy, October 26-30, 2015, Proceedings*, pages 185–200, 2015.
- [5] Benjamin Aminof, Aniello Murano, Sasha Rubin, and Florian Zuleger. Automatic verification of multi-agent systems in parameterised grid-environments. In *Proceedings of the 2016 International Conference on Autonomous Agents & Multiagent Systems (AAMAS 2016)*, pages 1190–1199, 2016.
- [6] Benjamin Aminof, Aniello Murano, Sasha Rubin, and Florian Zuleger. Prompt alternating-time epistemic logics. In *Principles of Knowledge Representation and Reasoning: Proc. of the 15th International Conference, (KR 2016)*, pages 258–267, 2016.
- [7] Benjamin Aminof and Sasha Rubin. First cycle games. *Information and Computation*, 2016.
- [8] Benjamin Aminof and Sasha Rubin. Model checking parameterised multi-token systems via the composition method. In *Proc. 8th International Joint Conference on Automated Reasoning, (IJCAR 2016)*, pages 499–515, 2016.
- [9] Benjamin Aminof, Sasha Rubin, Francesco Spegni, and Florian Zuleger. Liveness of parameterized timed networks. In *Automata, Languages, and Programming - 42nd International Colloquium, ICALP 2015, Kyoto, Japan, July 6-10, 2015, Proceedings, Part II*, pages 375–387, 2015.
- [10] Benjamin Aminof, Sasha Rubin, and Florian Zuleger. On the expressive power of communication primitives in parameterised systems. In *Logic for Programming, Artificial Intelligence, and Reasoning - 20th International Conference, LPAR-20 2015, Suva, Fiji, November 24-28, 2015, Proceedings*, pages 313–328, 2015.
- [11] Vince Bárány, Erich Grädel, and Sasha Rubin. Automata-based presentations of infinite structures. In Javier Esparza, Christian Michaux, and Charles Steinhorn, editors, *Finite and Algorithmic Model Theory*, pages 1–76. Cambridge University Press, 2011. Cambridge Books Online.
- [12] Francesco Belardinelli, Alessio Lomuscio, Aniello Murano, and Sasha Rubin. Verification of broadcasting multi-agent systems against an epistemic strategy logic. In *International Joint Conference on Artificial Intelligence (IJCAI 2017)*, 2017.
- [13] Francesco Belardinelli, Alessio Lomuscio, Aniello Murano, and Sasha Rubin. Verification of multi-agent systems with imperfect information and public actions. In *Proceedings of*

the 2017 International Conference on Autonomous Agents & Multiagent Systems (AAMAS 2017), 2017.

- [14] Raphael Berthon, Bastien Maubert, Aniello Murano, Sasha Rubin, and Moshe Vardi. Hierarchical strategic reasoning. In *IEEE Symposium on Logic in Computer Science (LICS 2017)*, 2017.
- [15] Roderick Bloem, Swen Jacobs, Ayrat Khalimov, Igor Konnov, Sasha Rubin, Helmut Veith, and Josef Widder. *Decidability of Parameterized Verification*. Synthesis Lectures on Distributed Computing Theory. Morgan & Claypool Publishers, 2015.
- [16] Roderick Bloem, Swen Jacobs, Ayrat Khalimov, Igor Konnov, Sasha Rubin, Helmut Veith, and Josef Widder. Decidability in parameterized verification. *SIGACT News*, 47(2):53–64, 2016.
- [17] Blai Bonet, Giuseppe De Giacomo, Hector Geffner, and Sasha Rubin. Generalized planning: Non-deterministic abstractions and trajectory constraints. In *International Joint Conference on Artificial Intelligence (IJCAI 2017)*, 2017.
- [18] Blai Bonet, Giuseppe De Giacomo, Hector Geffner, and Sasha Rubin. Generalized planning: Non-deterministic abstractions and trajectory constraints. In *ICAPS 2017 Workshop on Generalized Planning*, 2017.
- [19] Giuseppe De Giacomo, , Antonio Di Stasio, Aniello Murano, and Sasha Rubin. Imperfect information games and generalized planning. In *International Joint Conference on Artificial Intelligence (IJCAI 2016)*, 2016.
- [20] Hajime Ishihara, Bakhadyr Khoussainov, and Sasha Rubin. Some results on automatic structures. In *LICS 2002, 17th IEEE Symposium on Logic in Computer Science, 22-25 July 2002, Copenhagen, Denmark, Proceedings*, page 235, 2002.
- [21] Bakhadyr Khoussainov, André Nies, Sasha Rubin, and Frank Stephan. Automatic structures: Richness and limitations. In *LICS 2004, 19th IEEE Symposium on Logic in Computer Science, 14-17 July 2004, Turku, Finland, Proceedings*, pages 44–53, 2004.
- [22] Bakhadyr Khoussainov, André Nies, Sasha Rubin, and Frank Stephan. Automatic structures: Richness and limitations. *Logical Methods in Computer Science*, 3(2), 2007.
- [23] Bakhadyr Khoussainov, Sasha Rubin, and Frank Stephan. On automatic partial orders. In *LICS 2003, 18th IEEE Symposium on Logic in Computer Science, 22-25 June 2003, Ottawa, Canada, Proceedings*, pages 168–177, 2003.
- [24] Bakhadyr Khoussainov, Sasha Rubin, and Frank Stephan. Definability and regularity in automatic structures. In *STACS 2004, 21st Annual Symposium on Theoretical Aspects of Computer Science, Montpellier, France, March 25-27, 2004, Proceedings*, pages 440–451, 2004.
- [25] Bakhadyr Khoussainov, Sasha Rubin, and Frank Stephan. Automatic linear orders and trees. *ACM Transactions on Computational Logic*, 6(4):675–700, 2005.
- [26] Alex Kruckman, Sasha Rubin, John Sheridan, and Ben Zax. A myhill-nerode theorem for automata with advice. In *Proceedings Third International Symposium on Games, Automata, Logics and Formal Verification, GandALF 2012, Napoli, Italy, September 6-8, 2012.*, pages 238–246, 2012.
- [27] Aniello Murano, Giuseppe Perelli, and Sasha Rubin. Multi-agent path planning in known dynamic environments. In *PRIMA 2015: Principles and Practice of Multi-Agent Systems -*

18th International Conference, Bertinoro, Italy, October 26-30, 2015, Proceedings, pages 218–231, 2015.

- [28] Alexander Rabinovich and Sasha Rubin. Interpretations in trees with countably many branches. In *LICS 2012, Proceedings of the 27th Annual IEEE Symposium on Logic in Computer Science, Dubrovnik, Croatia, June 25-28, 2012*, pages 551–560, 2012.
- [29] Sasha Rubin. Automata presenting structures: A survey of the finite string case. *Bulletin of Symbolic Logic*, 14(2):169–209, 2008.
- [30] Sasha Rubin. Parameterised verification of autonomous mobile-agents in static but unknown environments. In *Proc. of the International Conference on Autonomous Agents and Multiagent Systems, (AAMAS 2015)*, pages 199–208, 2015.