



## ERC Advanced Grant 2017 Research proposal [Part B1]

# White-Box Self-Programming Mechanisms

## WHITEMECH

### Cover page:

- Principal investigator (PI): **Giuseppe De Giacomo**
- Host institution: **Università degli Studi di Roma “La Sapienza”**
- Proposal duration: **60 months**

We are witnessing an increasing availability of **mechanisms** that operate in nondeterministic (uncertain) environments and offer some form of programmability. These include manufacturing devices, smart objects and spaces, intelligent robots, dynamic business process management systems, and many others. All these mechanisms are currently being revolutionized by advancements in sensing (vision and language understanding) and actuation components (robot-arm movement and camera pointing). However, such mechanisms are held back by the fact that their logic is still based on hard-wired rules encoded in hand-crafted programs.

WHITEMECH aims at developing the science and the tools for a new generation of mechanisms to emerge: mechanisms that are able to **program themselves** without human intervention, automatically tailor their behavior so as to achieve desired goals, maintain themselves within safe boundaries in a changing environment, and follow regulations and conventions that evolve over time. Crucially, empowering mechanisms with self-programming carries significant risks and therefore we must be able to **balance power with safety**. For this reason WHITEMECH intends to realize mechanisms that are **white-box**, that is, guarded by human guided specifications and oversight, which, together with the synthesized behavior, are fully verifiable and comprehensible in human terms.

Remarkable recent discoveries by the applicant in **Reasoning about Action** and **Generalized Planning** in Artificial Intelligence, and their connections to Verification and Synthesis in Formal Methods, and Data-aware Processes in Databases, chart an unanticipated novel path to produce a breakthrough in realizing **powerful self-programming mechanisms**, while keeping them **human-comprehensible and safe by design**.

WHITEMECH will ground its scientific results upon three **driving applications**: smart manufacturing (Industry 4.0), smart spaces (IoT) and business process management systems (BPM).

## a Extended Synopsis of the scientific proposal

### LONG TERM VISION

Consider the following scenario.

*After a long week-end, the human supervisor inspects the manufacturing system and notices that a production line has slowed down significantly, though it is still producing. She queries the system on what it is doing. The system reveals to her the revised process, which is avoiding the use of the production island 176-176 by repurposing the tools in island 176-671 and sending items there. She then queries why the system has reprogrammed itself to do so. The system answers by showing that on Sunday 11:43pm the island 176-176 started to produce an unacceptable percentage of defective items, based on tests performed during production. So, the system restructured the process to achieve the specified objectives (quality and throughput) as it could in the presence of the faulty island, instead of shutting down the production line: the system analyzed the available capabilities and reprogrammed itself resulting in the current revised process.*

In the scenario above we have a **mechanism** (the manufacturing system) with multiple components, some of which possibly use Artificial Intelligence (AI) and in particular Machine Learning (ML), to provide sensing and acting capabilities (detecting defective items, reconfiguring tools parameters), that can be suitably organized to enact a *dynamic behaviour* (the production process) to meet its *specifications* (constraints on the product and production model). The mechanism has the ability to monitor and detect faulty parts of the process, by acquiring rich **relational data** about objects of interest, their properties and their relationships. Crucially, the mechanism has **self-programming** abilities that it can use to modify its current behavior, i.e., its logic, *without human intervention*. Notably, the mechanism can be *queried* to display, in terms understandable to humans, under which circumstances and for fulfilling which *specifications*, it reprogrammed itself. Moreover it can be queried on whether its *self-synthesized program* meets any dynamic property of interest to the human supervisor (e.g., for “what-if” analysis). In a slogan, the mechanism is **white-box**.

The overarching objective of WHITEMECH is to make this vision a reality:

WHITEMECH aims at laying the theoretical foundations and developing practical methodologies of a science and engineering of **white-box self-programming mechanisms**.

To make apparent the significance of this enterprise, WHITEMECH will ground its research in three driving applications considered of pivotal importance in the current socio-economic context, namely:

1. **Smart Manufacturing**, where significant research efforts are focusing on improving flexibility, agility and productivity of manufacturing systems, under the umbrella term *Industry 4.0*, or *4th industrial revolution*.<sup>1</sup>
2. **Internet of Things**, which is rising as a virtual fabric that connects “things” equipped with chips, sensors and actuators and allows for building **smart objects** and **smart spaces** with high levels of awareness of the environment and its human occupants.<sup>2</sup>
3. **Business Process Management**, which advocates explicit conceptual descriptions of a process to be enacted within an organization or possibly across organizations, and which is instrumental to business processes improvement, the top business strategy of CIOs according to Gartner.<sup>3</sup>

We observe that forms of self-programmability have been advocated in all the above contexts. For example, it is advocated that cyber-physical systems in Manufacturing or Internet of Things should be able to **adapt** their logic to deal with changes in their environment by **exploiting information gathered at runtime**. However, it is considered **impossible to determine a priori all possible adaptations** that may be needed; thus self-programming abilities would be highly desirable [114]. In Business Processes, it is considered important for the next generation of process management

<sup>1</sup>M. Lorenz et al. *Man and Machine in Industry 4.0: How Will Technology Transform the Industrial Workforce Through 2025?* The Boston Consulting Group. 2015.

<sup>2</sup>C. MacGillivray et al. *Worldwide Internet of Things Forecast Update, 2016-2020*. IDC. Doc # US40755516. 2016.

<sup>3</sup>Gartner Group. *BPM Survey Insights*. Gartner Report. <http://www.gartner.com/it/page.jsp?id=1740414>.

systems to allow processes to automatically **recover** when unanticipated exceptions occur, without explicitly defining a priori **recovery policies**, and **without the intervention of domain experts** at runtime [88]. These self-programming abilities would reduce costly and error-prone manual ad-hoc changes, and would relieve software engineers from mundane adaptation tasks [96]. Note that some of these concerns have been shared by *autonomic computing*, which has promoted self-configuration, self-healing, self-optimization, and self-protection, though by using policies provided by IT professionals [90]. Sophisticated languages and methodologies for streamlining the development of adaptation and exception handling have been devised [100, 126], however IT professionals still need to write the logic of these through **hard-wired rules encoded in hand-crafted programs** [34, 97, 81]. Although the interest is clearly apparent, currently self-programming abilities are missing or very limited in actual mechanisms operating in nondeterministic, uncertain, unpredictable environments.

WHITEMECH aims at **equipping such mechanisms with advanced self-programming abilities** so as, e.g., to enable them to:

- *Achieve desired goals*, that is guarantee that a certain desired state of affairs is eventually reached. In the above example a manufacturing system automatically reconfigures the fabrication process if some workstation is producing too many defective items, by changing the sequencing of processing units so as to temporarily cut-out the defective tool from the process, thus achieving an acceptable error rate.
- *Maintain themselves within a safe boundary* against unanticipated changes in the environment in which they operate. For example a smart space system may keep the desired temperature and humidity in a museum room at some desired level, even in presence of an unforeseen large crowd of visitors, possibly by momentarily repurposing other actuators, such as a secondary air conditioning system typically used only in case of failure of the main one.
- *Keep following regulations and conventions* that evolve over time while enacting their behavior. For example, to answer a new privacy regulation, a business process may refine its behavior to guarantee that the sensitive data are erased from the system before the completion of each process instance.

More generally, WHITEMECH wants to **enable mechanisms to act in an informed and intelligent way in their environment**, by changing the way they behave as a consequence of the information they acquire from the external world and exchange with the humans operating therein.

Since “*with great power comes great responsibility*”, introducing advanced forms of self-programming calls for the ability to make the behavior automatically synthesized by the mechanism **comprehensible** to human supervisors, who are thus able to control and guide it. So it is indeed crucial to develop self-programming mechanisms that **can be queried**, i.e., that are **white-box**: in every moment the mechanism can be queried for the status of its specifications, and on whether its behavior meets any dynamic property of interest to the human supervisors. Ultimately it is the fact that the resulting behavior is **comprehensible in human terms** that will make white-box self-programming mechanisms **trustworthy** [33, 99].

In the example of the first item above, both the reconfiguration goal (cutting out a defective tool) and how the fabrication process has been modified need to be explicitly understandable by the humans analyzing the manufacturing system. In the example of the second item, the sudden repurposing of the secondary air conditioning system also needs to be understandable to humans as a reaction to avoid violating certain requirements. Similarly, in the third item, the goal of erasing sensitive data from the system, and how this is achieved, must be understandable.

We stress that the need to move towards **white-box** approaches is advocated by a large part of the **AI community** [107] as well as the **CS community** [3], and has been recently taken up by DARPA within the context of machine learning, through the DARPA-BAA-16-53 “Explainable Artificial Intelligence (XAI)” program.<sup>4</sup> *Knowledge representation, the primary field of the PI, will be central for realizing the shift towards a white-box approach* [116].

WHITEMECH mechanisms will embrace a **model-based approach** to self-programming as opposed to a *model-blind (or function-based) approach* of some ML-techniques [39]. Clearly, there is a well justified enthusiasm for using such ML-techniques to develop smarter components. However, typically

<sup>4</sup><http://www.darpa.mil/program/explainable-artificial-intelligence>

these ML-components are black-boxes, in the sense that how they work remains opaque to humans [98, 117]. Though, there is currently much ongoing work on incorporating some forms of human-control into such ML-systems to provide safety guarantees [4, 120, 80]. WHITEMECH intends to support the integration of such ML-components with suitable safety guarantees into mechanisms.

## OBJECTIVES

Towards the goal of building **white-box self-programming mechanisms**, WHITEMECH will address 5 specific objectives, described below.

1. **Objective 1: Equip mechanisms with general self-programming abilities.** Mechanisms need general self-programming abilities, not restricted to a particular task, such as exception recovery, but ready to refine and modify the behavior of the mechanisms as new opportunities or constraints arise. Self-programming abilities are needed **while mechanisms are in operation**, that is, while mechanisms are executing, not at design time, as e.g., in [81]. WHITEMECH aims at developing advanced forms of **process synthesis** such as those studied in **Verification and Synthesis** in Formal Methods [103, 69] but that exploit recent advances in **AI Planning both Classical and Generalized** [76, 105, 119, 95, 54, 123, 61, 62, 32, 84, 79].
2. **Objective 2: Make self-programming mechanisms comprehensible and verifiable by humans.** WHITEMECH requires specifications, solutions (synthesized programs) and the relationship between them to be **comprehensible to humans**. This means that specifications and solution spaces must be **semantically** described at high-level, in human understandable terms, as advocated by **Knowledge Representation (KR)** [116, 11, 70, 71, 22, 127, 23, 93, 94]. Clearly, a crucial step towards comprehensibility is to be able to **verify** replanned behaviors against their specifications as recently advocated by **Reasoning about Action** in KR [106, 67, 36, 68, 121, 83, 38, 64, 7, 122, 48, 49, 112].
3. **Objective 3: Make self-programming mechanisms data-aware.** During the execution, new rich data about the world are observed, learned, or received as input. In particular, we are referring to **relational data** about objects of the world, their properties, and the relationships among them (as opposed to predefined continuous variables to be monitored as in Hybrid Systems [40, 102]). Self-programming mechanisms will leverage on semantic/intensional information to be able to interpret new data (extensional information) acquired, observed, and learned. Notice that this calls for a **first-order representation of the state**. New results on verifiability of **Data-Aware Processes** in the Databases and also BPM literature are crucial for this [2, 36, 65, 1, 8, 37, 89, 13, 48, 66, 87, 29].
4. **Objective 4: Support component-based approaches.** By no means should we consider programmable mechanisms to be composed of a single unit. Indeed, understandability calls for building **high-level components** that are relatively simple to understand, verify and combine. Then, it is crucial to study how **composing** “correct” components leads to an overall “correct” behavior. WHITEMECH will leverage work on composition and customization in Service Oriented Computing and more recently in **Reasoning about Action** in KR and **Generalized Planning** in AI [20, 118, 16, 57, 45], also looking at execution **monitoring** to detect failure, identify responsible components, and synthesize recovery [24, 41, 96].
5. **Objective 5: Integrate stochastic decisions and reinforcement learning.** Mechanisms may include components that have forms of decision making that resist definite analysis, which, e.g., make use of approximate models as in Markov Decision Processes (MDPs) and Reinforcement Learning (RL), possibly coupled with advanced deep learning techniques. However, while the actual execution in such components could be chosen stochastically, we do want to have **safety guarantees** on all possible generated executions [4, 120, 80]. We will study this objective in the context of **Generalized Planning and Synthesis under stochastic uncertainty** [76, 77, 35].

*Recent foundational results by the PI chart a novel path that within WHITEMECH will revolutionize **Reasoning about Action** in KR and **Generalized Planning** in AI by introducing rich objectives, data, and componentization in order to produce a **breakthrough in engineering self-programming mechanisms that are human-comprehensible and safe by design**.*

## METHODOLOGY AND SCIENTIFIC APPROACH

The scientific work in WHITEMECH will be methodologically structured into 3 broad research streams:

- **Principles and Foundations**, which will deal with the scientific foundations of white-box self-programming mechanisms.
- **Algorithms and Tools**, which will deal with the development of practical algorithms, optimizations and tools for realizing such mechanisms.
- **Applications and Evaluation**, which will evaluate white-box self-programming mechanisms in the three business critical applications above.

The **Principles and Foundations** and **Algorithms and Tools** streams will cut across 6 workpackages (WPs) roughly corresponding to the 5 objectives above. The **Applications and Evaluation** stream will be carried out in a specific WP, where the results of the project will be applied in the 3 driving applications domains. WPs are described in Part B2. Below we sketch WHITEMECH’s scientific approach, focusing on **novelty** and **feasibility**.

WHITEMECH will take from **Planning** in AI [76] (which in turn draws from **Reasoning about Action** in KR [106]) the idea of having **human-comprehensible** descriptions of the domain (the mechanism and its capabilities) and of the goals (the task specifications). However instead of considering simple domain specifications, as e.g., in PDDL, WHITEMECH will look at extensions that use **rich semantic descriptions** from KR and **componentization** from behavior compositions in Reasoning about Action and Generalized Planning, to which the PI has contributed significantly over the years [58, 46, 53, 110, 109, 52, 47, 50, 51, 10]. Particular attention will be given to computational effectiveness, in line with some recent exploratory work by the PI [48, 49, 29]. Notice that WHITEMECH does not aim at defining new concrete representation languages. Instead, it intends to use well-known formalisms such as BPMN, UML, OWL, etc. as concrete languages, though with a precise formal semantics to allow for automated reasoning, verification and synthesis, see e.g., [14, 56].

Another crucial difference with respect to Planning is that, instead of considering simple planning tasks, WHITEMECH will make use of the mathematically elegant theory of Reactive Synthesis [103] developed in **Verification and Synthesis** in the last 30 years [69], which however has not found widespread practical application because of the **intrinsic difficulties** of certain algorithms and constructions [124, 75]. We aim at **sidestepping these difficulties altogether**, by focusing on non-traditional kinds of specification formalisms developed recently in Reasoning about Action and Generalized Planning, such as LTL and LDL on finite traces, studied by the PI together with Moshe Vardi (Rice U, USA) [60, 61, 62], which are adopted in Planning in AI [9, 123, 32, 79, 113, 84] and in declarative business processes in BPM [125, 41, 54], as well as safe/co-safe LTL/LDL formulas, which have been shown to be more expressive than expected while remaining **well-behaved** [74, 73, 92, 72]. Solvers for these are substantially simpler than those for general reactive synthesis, being based on reachability and safety games, which are amenable to efficient implementations.

Moreover, like in Planning, but differently from Reactive Synthesis, we expect mechanisms to be able to handle quickly and efficiently most cases, i.e., those cases that do not require to solve difficult, “puzzle-like”, situations. Indeed, while the Planning community has concentrated on simpler forms of process synthesis (reachability of a state of affairs), it has developed a sort of **science of search algorithms for Planning**, which has brought about exceptional scalability improvements (orders of magnitude) in the last decade [105, 119, 95, 54]. WHITEMECH intends to devise new algorithms for solving reachability and safety games that are based on heuristic search as adopted in Planning, but also considering symbolic techniques adopted in synthesis by model checking [18]. The PI has been pioneering cross-fertilization of Planning and Synthesis since the end of the ’90 [59, 31, 111, 44, 101, 54]. More recently the PI has established **tight connections between Synthesis and Planning** [85, 86, 55, 19] as well as between Planning and Behavior Composition [108, 57, 45, 27].

Crucially, unlike current work in Planning and in Verification and Synthesis, which operate with a propositional representation of the state, WHITEMECH does not want to discard rich relational data, and hence it will need to consider a first-order representation of the state. This gives rise to infinite transition systems which are generally problematic to analyze. However the PI has already shown within the EU FP7-ICT-257593 ACSI: Artifact-Centric Service Interoperation, that such difficulties can be overcome in notable cases [15, 24, 8, 28] in the context of **Data-Aware Processes** in

Databases. A key discovery is that, under natural assumptions, these infinite-state transition systems admit faithful **abstractions** to finite-state ones, hence enabling the use of the large body of techniques developed within Verification and Synthesis in Formal Methods. Moreover, important advancements in understanding how to deal with such complexity have been established as well as relationships with Reasoning about Action in KR [82, 13, 30, 29, 10]. We will leverage on these ideas to lift our results so as to handle relational data.

WHITEMECH will be the spark that will bring together and cross-fertilize four distinct research areas, namely **Reasoning about Action** in Knowledge Representation, **Generalized Planning** in AI, **Verification and Synthesis** in Formal Methods, and **Data-aware Processes** in Databases.

*Given his track record, the PI is in a unique position to lead this cross-fertilization.*

The PI will dedicate 70% of his time to this project and reduce all his other duties accordingly (incl. teaching and other research projects). Apart from the PI and Sapienza personnel, the project will hire 4 Senior PostDocs, 6 Junior PostDocs, and 8 PhD students (cf. B2). Moreover, to foster cross-fertilization WHITEMECH will establish research exchanges with several research groups, including:

- **Reasoning about Action.** Collaborations with Yves Lesperance (York U., Toronto, Canada), Hector Levesque (U. Toronto, Canada), Sebastian Sardina (RMIT, Melbourne, Australia), and Yongmei Liu (Sun Yat-sen U., Guangzhou, China).
- **Generalized Planning.** Collaborations with Hector Geffner (UPF, Barcelona, Spain), Blai Bonet (U. Simon Bolivar, Caracas), and Alfonso Gerevini (U. Brescia, Italy).
- **Verification and Synthesis.** Collaborations with Moshe Vardi (Rice U.) on automata-based verification and synthesis, Sasha Rubin (U. Napoli, Italy) and Benjamin Aminof (TU Wien, Austria) on game-based verification and synthesis, and with Nello Murano (U. Napoli, Italy) on devising parallel synthesis algorithms to run on Multicore/GPUs [5, 6].
- **Data-Aware Processes.** Collaborations with Rick Hull (IBM Research, USA), Jianwen Su (UCSB, USA), Diego Calvanese and Marco Montali (U. Bolzano, Italy), and with Alessio Lomuscio (Imperial College, London, UK).

Moreover the PI will collaborate with Ronen Brafman (Ben-Gurion U.) to explore how to **incorporate components based on stochastic decision and RL** into mechanisms, and study temporally-extended goals in **non-Markovian** MDPs and RL (first ideas in [12, 92, 21]).

**Application and Evaluation.** WHITEMECH will ground its scientific results in the three real **applications domains** discussed above, which will provide continuous feedback to the scientific effort. Demonstrators will be developed to evaluate the practical effectiveness of the science and technology developed in the project. This work will also facilitate the take-up of WHITEMECH’s results in these more applied research areas. The PI and his group have contributed to these applications, see e.g., [42, 43, 63]. Moreover the PI has applied advanced science to real-cases in the past. In particular, he and his group have developed the Ontology-Based Data Access paradigm, possibly the most successful approach to Semantic Data Integration [25, 104, 26, 78, 17, 115, 91]. Recently, the PI founded the Sapienza Start-Up **OBDA Systems** (<http://www.obdasystems.com>) for its commercial exploitation.

## HIGH RISK, HIGH GAIN

WHITEMECH boldly aims at **bringing together and cross-fertilizing the above cited four distinct research areas** with overlapping interests but developed by different communities with different view-points, to produce a **breakthrough** that will make white-box self-programming mechanisms a reality. The **high risk** of this enterprise is **mitigated** by the PI expertise, who has already obtained initial foundational results, such as effective techniques for verification of data-aware processes [8, 48, 29], and feasibility results for synthesis, which sidestep some intrinsic difficulties of reactive synthesis algorithms and constructions (e.g., determinization) [60, 61, 62]. WHITEMECH will act as a **catalyst** for these areas and, together with current complementary advancements in ML, will bring about a **novel AI framework for self-programmability** that puts **human-comprehensibility at the center stage**.



## References

- [1] S. Abiteboul, P. Senellart, and V. Vianu. The ERC webdam on foundations of web data management. In *WWW*, pages 211–214, 2012. Work supported by the ERC Webdam (226513) project.
- [2] S. Abiteboul, V. Vianu, B. S. Fordham, and Y. Yesha. Relational transducers for electronic commerce. *J. Comput. Syst. Sci.*, 61(2):236–269, 2000.
- [3] ACM U.S. Public Policy Council and ACM Europe Policy Committee. Statement on algorithmic transparency and accountability. ACM, 2017.
- [4] D. Amodei, C. Olah, J. Steinhardt, P. Christiano, J. Schulman, and D. Mané. Concrete problems in AI safety. *CoRR*, abs/1606.06565, 2016.
- [5] R. Arcucci, U. Marotta, A. Murano, and L. Sorrentino. Parallel parity games: a multicore attractor for the Zielonka recursive algorithm. In *ICCS*, pages 525–534, 2017.
- [6] R. Arcucci, A. Murano, G. Perelli, and L. Sorrentino. A parallel model for reachability games. In *Submitted*, 2017.
- [7] J. Babb and J. Lee. Action language BC+: preliminary report. In *AAAI*, pages 1424–1430, 2015.
- [8] B. Bagheri Hariri, D. Calvanese, G. De Giacomo, A. Deutsch, and M. Montali. Verification of relational data-centric dynamic systems with external services. In *PODS*, pages 163–174, 2013.
- [9] J. A. Baier, F. Bacchus, and S. A. McIlraith. A heuristic search approach to planning with temporally extended preferences. *Artif. Intell.*, 173(5-6):593–618, 2009.
- [10] B. Banihashemi, G. De Giacomo, and Y. Lespérance. Abstraction in situation calculus action theories. In *AAAI*, pages 1048–1055, 2017.
- [11] C. Baral. *Knowledge Representation, Reasoning and Declarative Problem Solving*. Cambridge Univ. Press, 2010.
- [12] D. Beck and G. Lakemeyer. Reinforcement learning for Golog programs with first-order state-abstraction. *Logic Journal of the IGPL*, 20(5):909–942, 2012.
- [13] F. Belardinelli, A. Lomuscio, and F. Patrizi. Verification of agent-based artifact systems. *J. Artif. Intell. Res.*, 51:333–376, 2014.
- [14] D. Berardi, D. Calvanese, and G. De Giacomo. Reasoning on UML class diagrams. *Artif. Intell.*, 168(1-2):70–118, 2005.
- [15] D. Berardi, D. Calvanese, G. De Giacomo, R. Hull, and M. Mecella. Automatic composition of transition-based semantic web services with messaging. In *VLDB*, pages 613–624, 2005.
- [16] P. Bertoli, M. Pistore, and P. Traverso. Automated composition of web services via planning in asynchronous domains. *Artif. Intell.*, 174(3-4):316–361, 2010.
- [17] M. Bienvenu, P. Hansen, C. Lutz, and F. Wolter. First order-rewritability and containment of conjunctive queries in horn description logics. In *IJCAI*, pages 965–971, 2016. Work partially supported by the ERC CODA (647289) project.
- [18] R. Bloem, B. Jobstmann, N. Piterman, A. Pnueli, and Y. Sa’ar. Synthesis of reactive(1) designs. *J. Comput. Syst. Sci.*, 78(3), 2012.
- [19] B. Bonet, G. De Giacomo, H. Geffner, and S. Rubin. Generalized planning: Non-deterministic abstractions and trajectory constraints. In *IJCAI*, 2017.
- [20] A. Bouguettaya, Q. Z. Sheng, and F. Daniel, editors. *Web Services Foundations*. Springer, 2014.
- [21] R. Brafman, G. De Giacomo, M. Mecella, and S. Sardina. Service composition under probabilistic requirements. In *GenPlan Workshop at ICAPS*, 2017.

- [22] M. Brenner and B. Nebel. Continual planning and acting in dynamic multiagent environments. *Autonomous Agents and Multi-Agent Systems*, 19(3):297–331, 2009.
- [23] G. Brewka, S. Ellmauthaler, and J. Pührer. Multi-context systems for reactive reasoning in dynamic environments. In *ECAI*, pages 159–164, 2014.
- [24] D. Calvanese, G. De Giacomo, R. Hull, and J. Su. Artifact-centric workflow dominance. In *ICSOC*, pages 130–143, 2009.
- [25] D. Calvanese, G. De Giacomo, D. Lembo, M. Lenzerini, and R. Rosati. Tractable reasoning and efficient query answering in description logics: The *DL-Lite* family. *J. Autom. Reasoning*, 39(3):385–429, 2007.
- [26] D. Calvanese, G. De Giacomo, D. Lembo, M. Lenzerini, and R. Rosati. Data complexity of query answering in description logics. *Artif. Intell.*, 195:335–360, 2013.
- [27] D. Calvanese, G. De Giacomo, M. Lenzerini, and M. Y. Vardi. Regular open APIs. In *KR*, pages 329–338, 2016.
- [28] D. Calvanese, G. De Giacomo, M. Montali, and F. Patrizi. Verification and synthesis in description logic based dynamic systems. In *RR*, pages 50–64, 2013. Best paper award.
- [29] D. Calvanese, G. De Giacomo, M. Montali, and F. Patrizi. First-order  $\mu$ -calculus over generic transition systems and applications to the situation calculus. *Information & Computation*, 2017. To appear.
- [30] D. Calvanese, G. De Giacomo, and M. Soutchanski. On the undecidability of the situation calculus extended with description logic ontologies. In *IJCAI*, pages 2840–2846, 2015.
- [31] D. Calvanese, G. De Giacomo, and M. Y. Vardi. Reasoning about actions and planning in LTL action theories. In *KR*, pages 593–602, 2002.
- [32] A. Camacho, E. Triantafillou, C. J. Muise, J. A. Baier, and S. A. McIlraith. Non-deterministic planning with temporally extended goals: LTL over finite and infinite traces. In *AAAI*, pages 3716–3724, 2017.
- [33] C. Castelfranchi and R. Falcone. *Trust Theory: A Socio-Cognitive and Computational Model*. Wiley Series in Agent Technology. John Wiley & Sons, 2010.
- [34] V. G. Cerf. A long way to have come and still to go. *Commun. ACM*, 58(1):7, 2015.
- [35] K. Chatterjee and T. A. Henzinger. A survey of stochastic  $\omega$ -regular games. *J. Comput. Syst. Sci.*, 78(2):394–413, 2012.
- [36] J. Claßen and G. Lakemeyer. A logic for non-terminating Golog programs. In *KR*, pages 589–599, 2008.
- [37] J. Claßen, M. Liebenberg, G. Lakemeyer, and B. Zarriß. Exploring the boundaries of decidable verification of non-terminating Golog programs. In *AAAI*, pages 1012–1019, 2014.
- [38] J. Claßen, G. Röger, G. Lakemeyer, and B. Nebel. Platas - integrating planning and the action language Golog. *KI*, 26(1):61–67, 2012.
- [39] A. Darwiche. Human-level intelligence or animal-like abilities? *CoRR*, abs/1707.04327, 2017.
- [40] J. Davoren and A. Nerode. Logics for hybrid systems (invited paper). *Proceedings of the IEEE*, 88(7):985–1010, 2000.
- [41] G. De Giacomo, R. De Masellis, M. Grasso, F. M. Maggi, and M. Montali. Monitoring business metaconstraints based on LTL and LDL for finite traces. In *BPM*, pages 1–17, 2014.
- [42] G. De Giacomo, C. Di Ciccio, P. Felli, Y. Hu, and M. Mecella. Goal-based composition of stateful services for smart homes. In *OTM*, pages 194–211, 2012.
- [43] G. De Giacomo, M. Dumas, F. M. Maggi, and M. Montali. Declarative process modeling in BPMN. In *CAiSE*, pages 84–100, 2015.



- [44] G. De Giacomo, P. Felli, F. Patrizi, and S. Sardiña. Two-player game structures for generalized planning and agent composition. In *AAAI*, 2010.
- [45] G. De Giacomo, A. E. Gerevini, F. Patrizi, A. Saetti, and S. Sardiña. Agent planning programs. *Artif. Intell.*, 231:64–106, 2016.
- [46] G. De Giacomo, Y. Lespérance, and H. J. Levesque. ConGolog, a concurrent programming language based on the situation calculus. *Artif. Intell.*, 121(1-2):109–169, 2000.
- [47] G. De Giacomo, Y. Lespérance, and C. J. Muise. On supervising agents in situation-determined ConGolog. In *AAMAS*, pages 1031–1038, 2012.
- [48] G. De Giacomo, Y. Lespérance, and F. Patrizi. Bounded situation calculus action theories. *Artif. Intell.*, 237:172–203, 2016.
- [49] G. De Giacomo, Y. Lespérance, F. Patrizi, and S. Sardiña. Verifying ConGolog programs on bounded situation calculus theories. In *AAAI*, pages 950–956, 2016.
- [50] G. De Giacomo, Y. Lespérance, F. Patrizi, and S. Vassos. LTL verification of online executions with sensing in bounded situation calculus. In *ECAI*, pages 369–374, 2014.
- [51] G. De Giacomo, Y. Lespérance, F. Patrizi, and S. Vassos. Progression and verification of situation calculus agents with bounded beliefs. *Studia Logica*, 104(4):705–739, 2016.
- [52] G. De Giacomo, Y. Lespérance, and A. R. Pearce. Situation calculus based programs for representing and reasoning about game structures. In *KR*, 2010.
- [53] G. De Giacomo, H. J. Levesque, and S. Sardiña. Incremental execution of guarded theories. *ACM Trans. Comput. Log.*, 2(4):495–525, 2001.
- [54] G. De Giacomo, F. M. Maggi, A. Marrella, and F. Patrizi. On the disruptive effectiveness of automated planning for  $LTL_f$ -based trace alignment. In *AAAI*, pages 3555–3561, 2017.
- [55] G. De Giacomo, A. Murano, S. Rubin, and A. D. Stasio. Imperfect-information games and generalized planning. In *IJCAI*, pages 1037–1043, 2016.
- [56] G. De Giacomo, X. Oriol, M. Estañol, and E. Teniente. Linking data and BPMN processes to achieve executable models. In *CAISE*, pages 612–628, 2017.
- [57] G. De Giacomo, F. Patrizi, and S. Sardiña. Automatic behavior composition synthesis. *Artif. Intell.*, 196:106–142, 2013.
- [58] G. De Giacomo, R. Reiter, and M. Soutchanski. Execution monitoring of high-level robot programs. In *KR*, pages 453–465, 1998.
- [59] G. De Giacomo and M. Y. Vardi. Automata-theoretic approach to planning for temporally extended goals. In *ECP*, pages 226–238, 1999.
- [60] G. De Giacomo and M. Y. Vardi. Linear temporal logic and linear dynamic logic on finite traces. In *IJCAI*, pages 854–860, 2013.
- [61] G. De Giacomo and M. Y. Vardi. Synthesis for LTL and LDL on finite traces. In *IJCAI*, pages 1558–1564, 2015.
- [62] G. De Giacomo and M. Y. Vardi.  $LTL_f$  and  $LDL_f$  synthesis under partial observability. In *IJCAI*, pages 1044–1050, 2016.
- [63] L. de Silva, P. Felli, J. C. Chaplin, B. Logan, D. Sanderson, and S. M. Ratchev. Synthesising industry-standard manufacturing process controllers. In *AAMAS*, pages 1811–1813, 2017.
- [64] J. P. Delgrande and H. J. Levesque. A formal account of nondeterministic and failed actions. In *IJCAI*, pages 861–868, 2013.
- [65] A. Deutsch, R. Hull, F. Patrizi, and V. Vianu. Automatic verification of data-centric business processes. In *ICDT*, pages 252–267, 2009.

- [66] A. Deutsch, Y. Li, and V. Vianu. Verification of hierarchical artifact systems. In *PODS*, pages 179–194, 2016.
- [67] P. Doherty and J. Kvarnström. Talplanner: A temporal logic-based planner. *AI Magazine*, 22(3):95–102, 2001.
- [68] P. Doherty, J. Kvarnström, and F. Heintz. A temporal logic-based planning and execution monitoring framework for unmanned aircraft systems. *Autonomous Agents and Multi-Agent Systems*, 19(3):332–377, 2009.
- [69] R. Ehlers, S. Lafortune, S. Tripakis, and M. Y. Vardi. Supervisory control and reactive synthesis: a comparative introduction. *Discrete Event Dynamic Systems*, 27(2):209–260, 2017.
- [70] T. Eiter, E. Erdem, M. Fink, and J. Senko. Updating action domain descriptions. *Artif. Intell.*, 174(15):1172–1221, 2010.
- [71] T. Eiter, G. Gottlob, and T. Schwentick. The model checking problem for prefix classes of second-order logic: A survey. In *Fields of Logic and Computation, Essays Dedicated to Yuri Gurevich on the Occasion of His 70th Birthday*, pages 227–250, 2010.
- [72] P. Faymonville, B. Finkbeiner, M. N. Rabe, and L. Tentrup. Encodings of bounded synthesis. In *TACAS*, pages 354–370, 2017. Supported by the European Research Council (ERC) Grant OSARES (No. 683300).
- [73] E. Filiot, N. Jin, and J. Raskin. Antichains and compositional algorithms for LTL synthesis. *Formal Methods in System Design*, 39(3):261–296, 2011.
- [74] B. Finkbeiner and S. Schewe. Bounded synthesis. *STTT*, 15(5-6):519–539, 2013.
- [75] S. Fogarty, O. Kupferman, M. Y. Vardi, and T. Wilke. Profile trees for Büchi word automata, with application to determinization. In *GandALF*, 2013.
- [76] H. Geffner and B. Bonet. *A Concise Introduction to Models and Methods for Automated Planning*. Morgan & Claypool Publishers, 2013.
- [77] M. Ghallab, D. S. Nau, and P. Traverso. *Automated Planning and Acting*. Cambridge University Press, 2016.
- [78] G. Gottlob, G. Orsi, and A. Pieris. Query rewriting and optimization for ontological databases. *ACM Trans. Database Syst.*, 39(3):25:1–25:46, 2014. Work partially supported by the ERC DIADEM (246858) project.
- [79] J. Gutierrez, G. Perelli, and M. Wooldridge. Iterated games with LDL goals over finite traces. In *AAMAS*, pages 696–704, 2017. Work partially supported by the ERC RACE (291528) project.
- [80] D. Hadfield-Menell, A. Dragan, P. Abbeel, and S. Russell. The off-switch game. *CoRR*, abs/1611.08219, 2016.
- [81] D. Harel, G. Katz, R. Marelly, and A. Marron. First steps towards a wise development environment for behavioral models. *IJISMD*, 7(3):1–22, 2016.
- [82] B. B. Hariri, D. Calvanese, M. Montali, G. De Giacomo, R. De Masellis, and P. Felli. Description logic knowledge and action bases. *J. Artif. Intell. Res.*, 46:651–686, 2013.
- [83] S. Haufe, S. Schiffel, and M. Thielscher. Automated verification of state sequence invariants in general game playing. *Artif. Intell.*, 187:1–30, 2012.
- [84] K. He, M. Lahijanian, L. E. Kavrakı, and M. Y. Vardi. Reactive synthesis for finite tasks under resource constraints. In *IROS*, 2017.
- [85] Y. Hu and G. De Giacomo. Generalized planning: Synthesizing plans that work for multiple environments. In *IJCAI*, pages 918–923, 2011.
- [86] Y. Hu and G. De Giacomo. A generic technique for synthesizing bounded finite-state controllers. In *ICAPS*, 2013.

- [87] R. Hull, V. S. Batra, Y. Chen, A. Deutsch, T. Heath, and V. Vianu. Towards a shared ledger business collaboration language based on data-aware processes. In *ICSOC*, pages 18–36, 2016.
- [88] R. Hull, A. Koschmider, H. A. Reijers, and W. Wong. Fresh Approaches to Business Process Modeling (Dagstuhl Seminar 16191). *Dagstuhl Reports*, 6(5):1–30, 2016.
- [89] R. Hull, J. Su, and R. Vaculín. Data management perspectives on business process management. In *SIGMOD*, pages 943–948, 2013.
- [90] IBM. An architectural blueprint for autonomic computing. IBM White Paper, 2005.
- [91] E. Kharlamov, D. Bilidas, D. Hovland, E. Jimenez-Ruiz, D. Lanti, H. Lie, M. Rezk, M. Skjaeveland, A. Soyly, G. Xiao, D. Zheleznyakov, M. Giese, Y. Ioannidis, Y. Kotidis, M. Koubarakis, and A. Waaler. Ontology based data access in statoil. *Web Semantics: Science, Services and Agents on the World Wide Web*, 2017. In print.
- [92] B. Lacerda, D. Parker, and N. Hawes. Optimal policy generation for partially satisfiable co-safe LTL specifications. In *IJCAI*, pages 1587–1593, 2015.
- [93] H. J. Levesque. On our best behaviour. *Artif. Intell.*, 212:27–35, 2014.
- [94] H. J. Levesque. *Common Sense, the Turing Test, and the Quest for Real AI*. MIT Press, 2017.
- [95] N. Lipovetzky and H. Geffner. Best-first width search: Exploration and exploitation in classical planning. In *AAAI*, pages 3590–3596, 2017.
- [96] A. Marrella, M. Mecella, and S. Sardiña. Intelligent process adaptation in the smartpm system. *ACM TIST*, 8(2):25:1–25:43, 2017.
- [97] A. Marron, B. Arnon, A. Elyasaf, M. Gordon, G. Katz, H. Lapid, R. Marelly, D. Sherman, S. Szekely, G. Weiss, and D. Harel. Six (im)possible things before breakfast: Building-blocks and design-principles for wise computing. In *ACM/IEEE D&P@MoDELS*, pages 94–100, 2016.
- [98] V. Mnih, K. Kavukcuoglu, D. Silver, A. Graves, I. Antonoglou, D. Wierstra, and M. A. Riedmiller. Playing Atari with deep reinforcement learning. In *NIPS-2013 Workshop on Deep Learning*, 2013.
- [99] P. G. Neumann. Trustworthiness and truthfulness are essential. *Communication of ACM*, 60(8):26–28, June 2017.
- [100] M. Nidd, K. Bai, J. Hwang, M. Vukovic, and M. Tacci. Automated business application discovery. In *IM*, pages 794–797, 2015.
- [101] F. Patrizi, N. Lipovetzky, G. De Giacomo, and H. Geffner. Computing infinite plans for LTL goals using a classical planner. In *IJCAI*, pages 2003–2008, 2011.
- [102] A. Platzer. Logical analysis of hybrid systems: A complete answer to a complexity challenge. In *DCFS*, pages 43–49, 2012.
- [103] A. Pnueli and R. Rosner. On the synthesis of a reactive module. In *POPL*, pages 179–190, 1989.
- [104] A. Poggi, D. Lembo, D. Calvanese, G. De Giacomo, M. Lenzerini, and R. Rosati. Linking data to ontologies. *J. Data Semantics*, 10:133–173, 2008.
- [105] F. Pommerening, M. Helmert, and B. Bonet. Higher-dimensional potential heuristics for optimal classical planning. In *AAAI*, pages 3636–3643, 2017.
- [106] R. Reiter. *Knowledge in Action*. MIT Press, 2001.
- [107] S. Russell, D. Dewey, and M. Tegmark. Research priorities for robust and beneficial artificial intelligence. *AI Magazine*, 36(4), 2015.
- [108] S. Sardiña and G. De Giacomo. Realizing multiple autonomous agents through scheduling of shared devices. In *ICAPS*, pages 304–312, 2008.
- [109] S. Sardiña and G. De Giacomo. Composition of ConGolog programs. In *IJCAI*, pages 904–910, 2009.

- [110] S. Sardiña, G. De Giacomo, Y. Lespérance, and H. J. Levesque. On the semantics of deliberation in Indigolog - from theory to implementation. *Ann. Math. Artif. Intell.*, 41(2-4):259–299, 2004.
- [111] S. Sardiña, G. De Giacomo, Y. Lespérance, and H. J. Levesque. On the limits of planning over belief states under strict uncertainty. In *KR*, pages 463–471, 2006.
- [112] Z. G. Saribatur, C. Baral, and T. Eiter. Reactive maintenance policies over equalized states in dynamic environments. In *EPIA*, pages 709–723, 2017.
- [113] J. Segovia-Aguas, S. Jimenez, and A. Jonsson. Hierarchical finite state controllers for generalized planning. In *IJCAI*, 2017.
- [114] R. Seiger, S. Huber, and T. Schlegel. Toward an execution system for self-healing workflows in cyber-physical systems. *Software & Systems Modeling*, pages 1–22, 2016.
- [115] J. F. Sequeda and D. P. Miranker. A pay-as-you-go methodology for ontology-based data access. *IEEE Internet Computing*, 21(2):92–96, 2017.
- [116] Y. Shoham. Why knowledge representation matters. *Commun. ACM*, 59(1):47–49, 2016.
- [117] D. Silver, A. Huang, C. J. Maddison, A. Guez, L. Sifre, G. van den Driessche, J. Schrittwieser, I. Antonoglou, V. Panneershelvam, M. Lanctot, S. Dieleman, D. Grewe, J. Nham, N. Kalchbrenner, I. Sutskever, T. P. Lillicrap, M. Leach, K. Kavukcuoglu, T. Graepel, and D. Hassabis. Mastering the game of Go with deep neural networks and tree search. *Nature*, 529(7587):484–489, 2016.
- [118] S. Sohrabi, N. Prokoshyna, and S. McIlraith. Web service composition via the customization of Golog programs with user preferences. In *Conceptual Modeling: Foundations and Applications - Essays in Honor of John Mylopoulos*, pages 319–334, 2009.
- [119] M. Steinmetz and J. Hoffmann. State space search nogood learning: Online refinement of critical-path dead-end detectors in planning. *Artif. Intell.*, 245:1–37, 2017.
- [120] J. Taylor, E. Yudkowsky, P. LaVictoire, and A. Critch. Alignment for advanced machine learning systems. Technical report, Machine Intelligence Research Institute, 2017.
- [121] M. Thielscher. *Action Programming Languages*. Synthesis Lectures on Artificial Intelligence and Machine Learning. Morgan & Claypool Publishers, 2008.
- [122] S. T. To, T. C. Son, and E. Pontelli. A generic approach to planning in the presence of incomplete information: Theory and implementation. *Artif. Intell.*, 227:1–51, 2015.
- [123] J. Torres and J. Baier. Polynomial-time reformulations of LTL temporally extended goals into final-state goals. In *IJCAI*, pages 1696–1703, 2015.
- [124] M. Tsai, S. Fogarty, M. Y. Vardi, and Y. Tsay. State of Büchi complementation. *Logical Methods in Computer Science*, 10(4):1–27, 2014.
- [125] W. van der Aalst, M. Pesic, and H. Schonenberg. Declarative workflows: Balancing between flexibility and support. *Computer Science - R&D*, 23(2):99–113, 2009.
- [126] M. Vukovic, J. Hwang, J. J. Rofrano, and N. Anerousis. Blueshift: Automated application transformation to cloud native architectures. In *IM*, pages 778–792, 2017.
- [127] M. Wooldridge. *An Introduction to MultiAgent Systems*. John Wiley & Sons, 2nd edition, 2009.

## b Curriculum Vitae

### PERSONAL INFORMATION

Family Name, First Name: De Giacomo, Giuseppe  
 Researcher ID: [orcid.org/0000-0001-9680-7658](https://orcid.org/0000-0001-9680-7658)  
 Nationality: Italian  
 Date of birth: August 8, 1965  
 Home URL: <http://www.dis.uniroma1.it/~degiamoco>

### EDUCATION

1995 **PhD:** Dipartimento di Informatica e Sistemistica, Università degli Studi di Roma “La Sapienza”, Italy  
 Supervised by Prof. Maurizio Lenzerini.  
 1991 **Master:** Laurea in Ingegneria Elettronica (5 years), Facoltà di Ingegneria, Università degli Studi di Roma “La Sapienza”, Italy  
 1990–1991 **Erasmus Visiting Student** at University of Bristol, UK, with John W. Lloyd, preparing master thesis

### CURRENT POSITION

2006 – pres. **Full Professor** (Professore Ordinario), Università degli Studi di Roma “La Sapienza”, Italy

### PREVIOUS POSITIONS

2001 – 2006 **Associate Professor** (Professore Associato), Università degli Studi di Roma “La Sapienza”, Italy  
 1998 – 2001 **Assistant Professor** (Ricercatore), Università degli Studi di Roma “La Sapienza”, Italy  
 1996 – 1997 **Research Associate**, University of Toronto, Toronto, ON, Canada, working with Prof. Hector Levesque and Prof. Ray Reiter  
 1993/1994 **Visiting Scholar**, Stanford University, Stanford, CA, USA, working with Prof. Yoav Shoham

### FELLOSHIPS AND AWARDS

2016 **AAAI Fellow**, Fellow of the Association for the Advancement of Artificial Intelligence, Citation: *For significant contributions to the field of knowledge representation and reasoning, and applications to data integration, ontologies, planning, and process synthesis and verification.*  
 2015 **ACM Fellow**, Fellow of the Association for Computing Machinery, Citation: *For contributions to description logics, data management, and verification of data-driven processes*  
 2012 **ECCAI Fellow** (aka EurAI Fellow), Fellow of the European Association for Artificial Intelligence  
 2013 **Most influential paper in the decade** at the 11th International Conference on Service Oriented Computing (ICSOC’13) for the paper *Automatic Composition of E-services That Export Their Behavior* published at ICSOC03  
 2013 **Best paper award** at 7th international Conference on Web Reasoning and Rule Systems (RR’13) for the paper *Verification and Synthesis in Description Logic Based Dynamic Systems*  
 2013 **Miegunyah Distinguished Fellowship**, University of Melbourne, for the public lecture *Cognitive Robotics: The science of building intelligent autonomous robots and software agents*  
 2010 **IBM Open Collaborative Faculty Award 2010** for *Artifact-centric Business Process Modeling*, and **IBM Open Collaborative Faculty Award 2009** for *Radical Simplification of Artifact-Centric Business Process Modeling*

## SUPERVISION OF YOUNG RESEARCHERS

He supervised a number of **PhD students**, including Daniela Berardi (PhD in 2005, now working in industry), Fabio Patrizi (PhD in 2009, now assistant professor at Sapienza), Riccardo De Masellis (PhD in 2013, now PostDoc at FBK, Trento, Italy) and Paolo Felli (PhD in 2013, PostDoc at U. Nottingham), as well as **PostDocs**. He mentored a number of **early career researchers**, including Sebastian Sardina (now associate professor at RMIT, Melbourne), Marco Montali (now associate professor at U. Bolzano, Italy), Stavros Vassos (now funder of the StartUp Helvia.io), and Sasha Rubin (Marie Curie Fellow at U. Napoli, Italy).

## TEACHING ACTIVITIES

He has a wide teaching and academic experience. Since 1998, he has taught a large number of undergraduate and graduate courses at Sapienza. He has taught PhD courses in several European summer schools, including ESSLLI'03, ESSLLI'05, INFWEST'07. He has given several tutorials at top conferences, including ICSOC'04, WWW'05, ISWC'08, AAAI'10, IJCAI'15, IJCAI'16, ISWC'17.

## ORGANIZATION OF SCIENTIFIC MEETINGS

He regularly serves as PC member of many international conferences and workshops in AI and CS, including IJCAI ('09, '07, '05, '03, '01, '99, '95), AAAI ('15, '10, '07, '06, '04, '02, '00, '98), KR ('16 '12, '10, '08, '06, '02, '00, '98, '96), ECAI ('16, '10, '08), ICAPS ('15, '11, '09, '08, '07, '06), AAMAS ('13, '12, '02), PODS ('17\*, '15, '13, '11, '10, '09\*, '08, '07, '03, '02, '00 - \*PODS PC), ICDT ('15, '14, '11, '07, '05, '03, '99, '97), BPM ('12, '11, '10). Moreover he has been:

- Program Chair of KR'14
- Local organizer of KR'12
- Area chair of IJCAI'16, AAAI'12, IJCAI'11
- Senior PC of IJCAI'17, IJCAI'15, AAAI'16, AAAI'13, IJCAI'13, AAAI'11
- Workshops chair of AAAI'11, AAAI'10
- Organizer of 2003 Int. Workshop on Description Logics (DL'03) and of the 1st Cognitive Robotics Workshop at AAAI Fall Symp. (CogRob'98)

## INSTITUTIONAL RESPONSIBILITIES

- 2016 – curr. Member of the Department Directorate (Giunta di Dipartimento), Dipartimento di Ingegneria Informatica, Automatica e Gestionale, Sapienza
- 2011 - 2014 Director of the PhD Program in Computer Science and Engineering, Sapienza
- 2003 - 2008 Scientific Coordinator of the Erasmus Program at the Engineering School, Sapienza

## COMMISSIONS OF TRUST

- Artificial Intelligence J. (AIJ), Elsevier, Review Editor, since 2014; (Editorial Board Member, since 2013)
- Journal of Artificial Intelligence Research (JAIR), Associate Editor (2008 – 2015); (Editorial Board Member (2006 – 2008))
- Acta Informatica, Springer, Editorial Board Member, since 2015
- CoRR arXiv moderator for Artificial Intelligence, since 2014
- Vice President of Steering Committee of KR Inc., since 2016; in Steering Committee since 2014
- Member of W3C Recommendation Committee of OWL 2 Web Ontology Language Profiles (2012)
- Steering Committee of the Description Logics Workshop Series (2006 – 2009)
- AAAI Fellow Selection Committee, since 2017
- ECCAI/EurAI Fellow Selection Committee, since 2014

## MEMBERSHIPS OF SCIENTIFIC SOCIETIES

- Lifetime Member of Association for Computer Machinery (ACM)
- Lifetime Member of Association for Advancement of Artificial Intelligence (AAAI)
- Member of Italian Association for Artificial Intelligence (AI\*IA)



## Appendix: All on-going and submitted grants and funding of the PI (Funding ID)

### On-going grants

Project	Funding source	Amount	Period	Role of the PI	Relation to current ERC proposal
<b>ICE: Immersive Cognitive Environments</b>	Sapienza	40,000 EUR	2015-2018	PI	Focusses on smart spaces, on one of the application domains of WHITEMECH. It will be over when WHITEMECH starts.
<b>Theory and Techniques for Reasoning about Actions and High-level Agent Control in Multi-agent Domains;</b> NSFC Grant No. 61572535; Project PI: Yongmei Liu, Sun Yat-sen Univ. of Guangzhou	Natural Science Foundation of China	804,000 CNY (approx. 101,400 EUR)	2016-2019	Sapienza PI	Studies Reasoning about Action, but not in the context of self-programming. Funds visits to Sun Yat-sen Univ.
<b>SHF: Small: Pushing the Frontier of Linear-Time Model-Checking Technology;</b> NSF Grant No. 1319459; Project PI: Moshe Vardi, Rice University	NSF USA	304,582 USD (approx. 255,100 EUR)	2013-2018	Sapienza PI	Studies techniques for model checking. Funds visits to Rice. It will be over when WHITEMECH starts.

### Grant applications

Project	Funding source	Amount	Period	Role of the PI	Relation to current ERC proposal
<b>Beyond One-shot Planning: Cognitive Agent Planning Programs;</b> ARC ID DP180102600; Project PI: Sebastian Sardina, RMIT, Australia	ARC Grant Australia	392,541 AUD (approx. 261,200 EUR)	2017-2020	Sapienza PI	Studies composition. Funds visits to RMIT.

### c Ten-year track record

PI's research concerns theoretical, methodological and applicative aspects in different areas of AI and CS, including: LTL and LDL over finite traces; Bounded situation calculus; Decidability of data-aware processes; Generalized planning by model checking and automata-theoretic techniques from Verification; ConGolog programming language based on Situation Calculus; Ontology Based Data Access (OBDA); DL-lite family: description logics with tractable data complexity; Reasoning on UML Class Diagrams; Service composition and synthesis; Regular path queries for view-based query answering in graph databases; Data integration with description logics constraints; Conjunctive query answering in description logics; Correspondence between description logics and logics of programs. This work has deeply impacted several areas of AI and CS.

### IMPACT MEASURES

The PI is the author of more than 250 publications in top journals and conferences, including the following CORE A\*/A conferences and journals (since 2007 in bold, career in italic): IJCAI (**17**, *22*, A\*), AAAI (**11**, *16*, A\*), KR (**10**, *20*, A\*), AAMAS (**7**, *8*, A\*), ICAPS (**5**, *5*, A\*), PODS (**3**, *9*, A\*), VLDB (**2**, *3*, A\*), ISWC (**2**, *2*, A), CAiSE (**2**, *4*, A), ECAI (**3**, *4*, A), BPM (**3**, *3*, A), ICSOC (**1**, *4*, A), ICDT (**1**, *3*, A), Artif. Intell. (**4**, *7*, A\*), J. Comput. Syst. Sci. (**2**, *3*, A\*), Inf. Syst. (**1**, *2*, A\*), J. Log. Comput. (**1**, *4*, A), J. Artif. Intell. Res. (**1**, *2*, A), J. Autom. Reasoning (**1**, *1*, A), ACM Trans. Comput. Log. (**1**, *2*, A), Theor. Comput. Sci. (**1**, *2*, A). A comprehensive list of publication can be found on DBLP <http://dblp.uni-trier.de>.

According to Google Scholar, August 2017, the PI's **h-index is 68**, with **17164 citations** and **his i10-index is 193**. These values are among the highest in Europe in AI and CS and make the PI the 3rd most cited CS author in Italy, according to a study available at <http://via-academy.org>. Currently the PI's 5 most cited papers are (citations from Google Scholar. August 2017):

1. **Tractable reasoning and efficient query answering in description logics: The DL-Lite family.** D. Calvanese, G. De Giacomo, D. Lembo, M. Lenzerini, R. Rosati. *Journal of Automated Reasoning* 39 (3), 385-429, 2007 - **1174** cit.
2. **ConGolog, a concurrent programming language based on the situation calculus.** G. De Giacomo, Y. Lesperance, H. Levesque. *Artif. Intell.* 121(1-2): 109-169 (2000) - **659** cit.
3. **Linking data to ontologies.** A. Poggi, D. Lembo, D. Calvanese, G. De Giacomo, M. Lenzerini, R. Rosati. *Journal on data semantics X*, 133-173, 2008 - **578** cit.
4. **Reasoning on UML class diagrams.** D. Berardi, D. Calvanese, G. De Giacomo. *Artif. Intell.* 168(1-2): 70-118 (2005) - **557** cit.
5. **Automatic Composition of E-services That Export Their Behavior.** D. Berardi, D. Calvanese, G. De Giacomo, M. Lenzerini, M. Mecella. *ICSOC 2003*: 43-58 - **520** cit. – Awarded as “the most influential ICSOC paper in the last 10 years” at ICSOC 2013.

### REPRESENTATIVE PUBLICATIONS (MOST RELEVANT FOR PROJECT)

1. **Generalized Planning: Non-Deterministic Abstractions and Trajectory Constraints.** B. Bonet, G. De Giacomo, H. Geffner, S. Rubin. *IJCAI 2017* - CORE A\*.
2. **Agent planning programs.** G. De Giacomo, A. Gerevini, F. Patrizi, A. Saetti, S. Sardina. *Artif. Intell.* 231: 64-106 (2016) - CORE A\*.
3. **Bounded situation calculus action theories.** G. De Giacomo, Y. Lesperance, F. Patrizi. *Artif. Intell.* 237: 172-203 (2016) - CORE A\*.
4. **Regular Open APIs.** D. Calvanese, G. De Giacomo, M. Lenzerini, M. Vardi. *KR 2016*: 329-338. - CORE A\*.
5. **Synthesis for LTL and LDL on Finite Traces.** G. De Giacomo, M. Vardi. *IJCAI 2015*: 1558-1564 - CORE A\*.
6. **Data complexity of query answering in description logics.** D. Calvanese, G. De Giacomo, D. Lembo, M. Lenzerini, R. Rosati. *Artif. Intell.* 195: 335-360 (2013) - CORE A\*.
7. **Automatic behavior composition synthesis.** G. De Giacomo, F. Patrizi, S. Sardina. *Artif. Intell.* 196: 106-142 (2013) - CORE A\*.
8. **Verification of relational data-centric dynamic systems with external services.** B. Bagheri Hariri, D. Calvanese, G. De Giacomo, A. Deutsch, M. Montali. *PODS 2013*: 163-174 - CORE A\*.

9. **On simplification of schema mappings.** D. Calvanese, G. De Giacomo, M. Lenzerini, M. Vardi. *J. Comput. Syst. Sci.* 79(6): 816-834 (2013) - CORE A\*.
10. **View-based query answering in Description Logics: Semantics and complexity.** D. Calvanese, G. De Giacomo, M. Lenzerini, R. Rosati. *J. Comput. Syst. Sci.* 78(1): 26-46 (2012) - CORE A\*.

### INVITED TALKS (SELECTION)

- On **LTL and LDL on Finite Traces: Reasoning, Verification, and Synthesis:** GenPlan@ICAPS 2017 (Keynote), Pittsburgh, USA; SR 2016 (Keynote), New York, USA; Highlights of Logic, Games and Automata 2015 (Keynote), Prague, CZ; ICAPS 2013 (Keynote), Rome, IT.
- On **AI Foundations for Data-Aware BPM:** CBPM@BPM 2017, Barcelona, SP, (Keynote); York U., (Distinguished Lassonde Lecture), Canada, 2017; U. of Toronto, Canada, 2017; WS-FM:FASOCC@BPM 2014, Eindhoven, NL; ECAI 2014 (Keynote for Frontiers of AI), Prague, CZ.
- On **Bounded Situation Calculus:** WS. Formal Methods in AI 2017 (Keynote), U. “Federico II”, Naples, IT; TIME 2015 (Keynote), Kassel, GE; WS. HYBRIS 2015 (Keynote), Potsdam, GE.
- On **Service Composition and Synthesis:** ICSOC 2013 (Invited talk for prize as most influential paper of decade), Berlin, GE; U. Brescia, IT 2012; U. of Toronto Canada, 2010; York University, Toronto, Canada, 2010; INFINT WS 2009, Bertinoro, IT; MSI 2005 Berlin, GE, Caen, FR.
- On **Ontology-based Data Access and Integration:** DL 2013 (Keynote), Ulm, Germany; U. of Toronto, Canada, 2010; Semantic Days Conference 2009, Stavanger, NO; IBM Research Center Watson, Hawthorne, NY, USA 2008.

### RESEARCH PROJECTS (SELECTION)

The PI has leaded several projects, including: (2005–2008) EU FP6-7603 **TONES: Thinking OntologiES**, PI for Sapienza, value: EUR 264,000 (total value: EUR 1,438,910), from final review: *The TONES project can be considered as a success story of a FET-project in terms of scientific achievements*; (2010–2013) EU FP7-ICT-257593 **ACSI: Artifact-Centric Service Interoperation**, PI for Sapienza, and Scientific Coordinator for the whole project, value: EUR 435,000 (total value: EUR 3,243,937), from final review: *The scientific productivity of ACSI has been extraordinary, leading to a great impact on BPM, DB and AI research, as indicated by invited talks and tutorials on ACSI results in many high profile conferences, and by the addition of ACSI topics to the list of conference and workshop topics in these areas.*; (2012–2016) EU FP7-IST-IP-318338 **OPTIQUE: Scalable End-user Access to Big Data**, key personnel of Sapienza, value: EUR 802 488 (total value: EUR 9,838,329); (2009–2014) Open Collaboration Research W0954341 with Rick Hull of IBM T. J. Watson Research Center, NY, on **Data Aware Business Processes and Operation, through An Artifact-Centric Approach**, PI, value: USD 45,000; (2010–2012) UK Royal Society International Joint Project 2009/R2 on **Web Services Automatic Synthesis through ATL Symbolic Model Checking**, with Alessio Lomuscio, Imperial College London, total value: GBP 12,000; (2012–2014) Australian Research Council (ARC) Discovery Project DP120100332 **Optimisation of Embedded Virtual Complex Systems by Re-using a Library of available component**, with Sebastian Sardina of RMIT and Maurice Pagnucco of Univ. of Sidney, PI for Sapienza, total value: AUD 406,278; (2013–2015) Sapienza Ateneo Project **Spiritiles: Spiritlet-based Smart Spaces**, PI, value: EUR 60,000. Currently active project are listed in the CV.

### INDUSTRIAL LEADERSHIP

A pioneer project in 2003 with IBM Tivoli Lab highlighted several fundamental limitations in using semantic technologies available at that time for knowledge management. This gave rise to a novel research line that led to a new kind of Description Logic, called DL-Lite, and a new paradigm for accessing data through semantic technologies, called Ontology-based Data Access (OBDA). Since then, the PI led several industrial explorations promoting OBDA for data integration, preparation and discovery, which involved private and public organizations, such as the Monte dei Paschi di Siena Bank, Italian Ministry of Finance and Economics, Telecom Italia, Bloomberg, Italian Automobil Club (current). The success of such explorations has led the PI to constitute the Sapienza Start Up **OBDA Systems** (<http://www.obdasystems.com>) in Feb. 2017. The PI has been also a Contributor to W3C recommendation of OWL 2 Web Ontology Language Profiles (<https://www.w3.org/TR/owl2-profiles/>).